




nordic summit

Billund | 21st-22nd September 2026

You Don't Need to Know Everything to Start

A Real Path Into Data with Microsoft Fabric



Lun@Lake



Fellowwind



Gold Sponsors

oneflow



Silver Sponsors



502nm



ColorCloud



Bronze Sponsors



nordic summit

Sandra Kiel

Wants to start working with Microsoft Fabric

- Co-Founder & Managing Director @ Event Punks, Microsoft MVP
- Focus Topics: Gamification, Microsoft 365 & Power Platform
- Fun Fact: My natural response to a new challenge is usually: "Let's turn it into a game."





nordic summit


Hanna Schwab

Works with Microsoft Fabric every day

- Data & AI Lead @ tecle group, Microsoft Foundry MVP
- Focus Topics: Agentic AI, process automation
- Fun Fact: 5 years ago I didn't know what Azure was (oops)



Session roadmap



The Confidence Gap

Anyone who already knows Power BI, Copilot, or Power Apps well often hits a higher bar when starting something completely new.

The reason: for the first time in a while, that expert status resets to almost zero.

What makes learning a new skill feel hard?

- **Too Many Options/Features**

The landscape feels overwhelming before you've even started.

- **No Clear Starting Point**


Everyone else already seems to know where to begin.

- **Fear Of Basic Questions**


Asking about a Lakehouse feels uncomfortable when everyone else seems like an expert.

- **Tried Once, Got Confused**

Started once, hit a wall and "will continue later".



Fabric "Mental Model"



- Every workload above reads and writes the same stack of data in OneLake, which is why the export/import step between different tools disappears

Is Fabric Even Right For Me?



Fabric Fits When...

- Multiple roles (engineers, analysts, ...) need one shared platform
- Multiple data sources need to be combined to build one shared source of truth
- You're already in the Microsoft ecosystem
- You want out of tool sprawl: separate ADF, Synapse, Power BI, Databricks, ...

Fabric Doesn't Fit (Yet) When...

- The use case is small, Excel/PowerBI genuinely covers it
- You're strategically multi-cloud
- You have legacy system constraints
- The team has no Microsoft skillset, and switching doesn't pay off (yet)



Where Fabric Shows Up In Daily Work

3 use cases for Microsoft Fabric in our daily work life:

1

The Monthly Power BI Refresh

A 40-person sales team where every regional manager rebuilds the same pipeline report slightly differently in Excel every Monday.



2

The One-Person Pipeline

A single biotech data scientist merging 6 messy lab-instrument CSV feeds by hand in Python every day, 3 hours a day, nobody else touches this data.



3

The Nightly VLOOKUP

A 60-person retail chain manually combining POS and inventory every night for one weekly leadership report



4


The Copilot Studio Agent

An agent that answers live from OneLake, instead of from an export that was already stale the moment it was created



- In all three cases, the tool on the surface doesn't change, only where the data behind it comes from.

Before Fabric vs. With Fabric



Before Fabric vs. With Fabric



Business Understanding

Data & Meaning & Action

Decision-making

Agility & Ownership

TODAY

Depen

Fragm

Insula

Developed by specialized engineers



WITH FABRIC

AI

Shared Language for Teams and AI

Consolidation into a unified data model

Company-wide data foundation in OneLake

Single source of truth with business context

Single source of truth with business context

Transparent & centrally managed

Transparent & centrally managed

So it's basically the same way of working with tools like Power BI but with a shared data layer underneath it that wasn't there before!?

Fabric "Mental Model"

You don't need to know & use every feature right from the start!



Tenant-level

Tenant

Capacity-level

Capacity 1

Capacity n

Workspace-level

Workspace 1

Workspace 2

Workspace n

Item-level

Data Warehouse

Lakehouse

Semantic Model

Object-level

Dbo.Customer

VM.MySQLView

What's a Microsoft Fabric Capacity?!

- Capacity = compute power you purchase to run your workloads
- Size of capacity = amount of computation power available
- Billing: pay-as-you-go via Azure (hourly, pausable) or reserved (cheaper, annual commitment)
- Measured in **Capacity Units (CUs)**, size doubles per tier: F2, F4, F8, F16 ... up to F2048
- Capacity is not a fixed monthly quota (CU consumption depends on average load, prolonged overuse, and backlog accumulation)

Sample Calculation (Estimates)*

Data Warehouse Query (1 GB)	~10 CU-Sek.
Spark ETL-Job (1 GB)	~50 CU-Sek.
Power BI Report-Interaktion	~5–20 CU-Sek.
OneLake-Lesezugriff (16 MB)	~4 CU-Sek.

Example:

- 500 GB/day ingestion via Dataflow Gen2 $\approx$ 50,000 CU-seconds.
- F8 capacity provides 8 CU/s on a sustained basis
- If we would compare this to the computing power theoretically available over 24 hours, this workload would account for about **7% of it**

Fabric "Mental Model"

You don't need to know & use every feature right from the start!



Tenant-level

Tenant

Capacity-level

Capacity 1

Capacity n

Workspace-level

Workspace 1

Workspace 2

Workspace n

Item-level

Data Warehouse

Lakehouse

Semantic Model

Object-level

Dbo.Customer

VM.MySQLView

Fabric "Mental Model"

You don't need to know & use every feature right from the start!



Data exchange within Fabric can take place at one of three levels:

Workspace-level

Workspace 1

Workspace 2

Workspace n

Item-level

Data Warehouse

Lakehouse

Semantic Model

Object-level

Dbo.Customer

VM.MySQLView

Lakehouse vs. Warehouse



Lakehouse

- Files and tables together in one place
- Spark and notebooks, flexible schema
- Engineering and data-science leaning




Warehouse

- Pure T-SQL, strict schema
- Guaranteed ACID transactions
- Classic BI and analyst teams

No Code or Pro Code?


Warehouse Requirements?

Complexity of data?




- The choice follows the data needs → In medallion architecture you'd typically use Lakehouses for Bronze/Silver and Warehouses for Gold layers

But where does my data come from?



Best case vs good case: Medallion Architecture in Microsoft Fabric



- Even small projects benefit from at least a light Bronze/Silver split, instead of jumping straight to Gold!

Live Demo: End-to-End Data Analytics

Business questions answered from a common data estate: from raw data to a published report or data agent

1

Integration

Pipeline or Shortcut into the Lakehouse (Bronze layer)

2

Transformation

Dataflow Gen2, proof that low-code is enough to start (Silver layer)

3

Insights

Power BI report or Fabric Data agents to answer your business questions (Gold / Access layer)

Copilot Is Already Everywhere in Fabric

Not a separate product, it's built into the workloads you just saw. Needs F2+ or P SKU (not available on trial)



Data Engineering & Science

Notebook code generation and "Fix with Copilot" for errors 😊



Data Factory

Natural language turns into pipeline logic and transformations – so efficient!



Data Warehouse & SQL Database

Natural language to T-SQL, query completion, execution plan analysis and more



Power BI & Real-Time Intelligence

Auto-generated reports, natural language to KQL for real-time dashboards – hell yes, such a time-saver!



- In every one of these workloads, Copilot does the same job: natural language becomes runnable code or a report directly.

Copilot vs. Fabric: When to Use What



- For **conversational AI** & automations
- **Direct access** to existing systems (M365, Dataverse, APIs)
- Ideal for FAQ-Bots, process automations, **simple integrations**
- **No replacement** for central data storage or BI



- For **data integration, governance & analytics**
- Unified **Single Source of Truth** (OneLake, Semantic Models)
- Base for **BI, AI-training, RAG** & compliance
- Scalable for **Big Data & complex use cases**



- **Fabric provides curated data & governance**
- **Copilot uses them for intelligent chats**
- Example: Sales-Assistant → KPIs & documents in Fabric as knowledge base for intelligent Chatbot in natural language

But where Do I Actually Start?

What You Don't Need on Day One

No Spark required to get started

No new security model needed upfront

No CI/CD required before your first project

Where You Start Instead

1

Pick One Thing

Don't migrate your whole data estate, chose one simple use case for your first MS Fabric project

2


Land It in OneLake

Your entry point; Trial or F2 capacity is enough to learn end to end.

3

Use What You Know

Dataflow Gen2 first, that's like Power Query. Publish a report, try Fabric Data Agents and try new features over time





Q&A



Lun@Lake



Fellowwind



Gold Sponsors

oneflow



Silver Sponsors



502nm



ColorCloud



Bronze Sponsors



THANK YOU